

#GPU

```
import tensorflow as tf
from tensorflow.keras import layers, models
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from sklearn.model_selection import train_test_split
```

```
# Set up parameters
AUTO = tf.data.AUTOTUNE
IMAGE_SIZE = [224, 224] # adjust size as needed
batch_size = 32
gcs_pattern_train = 'gs://grad_proj/train/**/*.jpg'
gcs_pattern_test = 'gs://grad_proj/test/**/*.jpg'
num_images_per_class = 100
num_images_per_class = 100
```

```
# Define data loading functions
def parse_image(filename, label):
    img = tf.io.read_file(filename)
    img = tf.image.decode_jpeg(img, channels=3)
    img = tf.image.resize(img, IMAGE_SIZE)
    img = tf.cast(img, tf.float32) / 255.0
    return img, label
```

```
# Define data loading functions
def load_dataset(gcs_pattern, num_images):
    filenames = tf.io.gfile.glob(gcs_pattern)
    random.shuffle(filenames) # Shuffle the filenames for randomness
    filenames = filenames[:num_images] # Select a subset of filenames
    labels = [1 if 'dog' in filename else 0 for filename in filenames]
    dataset = tf.data.Dataset.from_tensor_slices((filenames, labels))
    dataset = dataset.map(parse_image, num_parallel_calls=AUTO)
    return dataset
```

```
import random
# Load datasets
train_dataset = load_dataset(gcs_pattern_train, num_images_per_class)
test_dataset = load_dataset(gcs_pattern_test, num_images_per_class)

# Assuming your train dataset is named 'train_dataset'
train_dataset_length = tf.data.experimental.cardinality(train_dataset)

print(f"Train dataset length: {train_dataset_length}")
```

Train dataset length: 100

Resources X

...

You are not subscribed. [Learn more.](#)

You currently have zero compute units available. Resources offered free of charge are not guaranteed. Purchase more units [here](#).

[Manage sessions](#)

Want more memory and disk space? X

[Upgrade to Colab Pro](#)

Python 3 Google Compute Engine
backend (GPU)

Showing resources from 03:17 to 03:34

System RAM
4.2 / 12.7 GB



GPU RAM
2.8 / 15.0 GB



Disk
26.9 / 78.2 GB



```
# Define and compile the model
model = models.Sequential()
model.add(layers.Conv2D(32, kernel_size=(3, 3), padding='same'))
model.add(layers.BatchNormalization())
model.add(layers.MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(layers.Conv2D(64, kernel_size=(3, 3), padding='same'))
model.add(layers.BatchNormalization())
model.add(layers.MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(layers.Conv2D(128, kernel_size=(3, 3), padding='same'))
model.add(layers.BatchNormalization())
model.add(layers.MaxPooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(layers.Flatten())
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dropout(0.1))
model.add(layers.Dense(64, activation='relu'))
model.add(layers.Dropout(0.1))
model.add(layers.Dense(1, activation='sigmoid'))

model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
```

```
import time
```

```
start_time = time.time()
```

```
history = model.fit(
    train_dataset.shuffle(1000).batch(batch_size).prefetch(1),
    epochs=10,
    validation_data=test_dataset.batch(batch_size).prefetch(1),
)
```

```
end_time = time.time()
```

```
elapsed_time = end_time - start_time
```

```
print(f"Training took {elapsed_time} seconds.")
```

```
Epoch 1/10
4/4 [=====] - 29s 3s/step -
Epoch 2/10
4/4 [=====] - 12s 2s/step -
Epoch 3/10
4/4 [=====] - 16s 3s/step -
Epoch 4/10
4/4 [=====] - 11s 2s/step -
Epoch 5/10
4/4 [=====] - 16s 3s/step -
Epoch 6/10
4/4 [=====] - 16s 3s/step -
Epoch 7/10
4/4 [=====] - 10s 2s/step -
Epoch 8/10
4/4 [=====] - 15s 3s/step -
Epoch 9/10
4/4 [=====] - 15s 3s/step -
Epoch 10/10
4/4 [=====] - 10s 2s/step -
Training took 161.66323280334473 seconds.
```

[Change runtime type](#)